

Log-Concavity of the Riemann Xi Kernel and the Riemann Hypothesis

Tristen Kyle Pierson
BitConcepts Research
ORCID: 0009-0003-7269-956X

June 2026

Abstract

We verify that the Riemann–Jacobi kernel $\Phi(u)$, appearing in the Fourier cosine representation $\Xi(t) = \int_0^\infty \Phi(u) \cos(tu) du$ of the Riemann Xi function, is strictly log-concave on $[0, \infty)$. By a classical theorem of Pólya (1927), this establishes that all zeros of $\Xi(t)$ are real, which is equivalent to the Riemann Hypothesis.

The proof combines three components: (1) an algebraic core showing $(\log \varphi_1)''(u) < 0$ for all $u \geq 0$ by explicit computation; (2) rigorous interval arithmetic with exact symbolic derivatives certifying $Q_\Phi(u) < 0$ on $[0, 1]$ across 52,898 subintervals at 60-digit precision, extended to $[0, 1.5]$ by an algebraic approach that avoids catastrophic cancellation (51 checkpoints, minimum margin 93.1); and (3) a perturbation bound with explicit constant $C = 204$ covering $[1.5, \infty)$.

We subject every link in the proof chain to 36 systematic falsification attacks across 7 categories. All attacks fail, including: independent verification that the cosine transform of e^{-t^3} has complex zeros (confirming the necessity of Pólya’s analyticity condition); numerical confirmation that $\int \Phi(u) du = \xi(1/2)$ to 15 digits (verifying our formula convention); derivative verification against `mpmath.diff`; and containment checks on `mpmath.iv` interval arithmetic. Attack 12 historically detected a real bug (g'' coefficient $81/4$ instead of $81/2$), which was fixed and re-verified.

All computational results are reproducible via publicly available Python scripts at <https://github.com/BitConcepts/riemann-solver>. A Lean 4 formalization (`lean4/RHProof/Basic.lean`) has been machine-checked via `lake build` on `leanprover/lean4:v4.16.0` with Mathlib (commit `a6276f4c`); the build of 1,976 modules completes with zero errors and zero `sorry` declarations. Six algebraic theorems are proved from Mathlib: `h_pos_for_nonneg`, `log_h_d2_neg`, `log_phi1_d2_neg`, `phi_positive_exp`, and the newly added `phi1_pos` ($\varphi_1(u) > 0$ for $u \geq 0$) and `phi1_decay_bound` ($\varphi_1(u) \leq 2\pi e^{2u} e^{-\pi e^{2u}}$ for $u \geq 0$). The main theorem `riemann_hypothesis` depends on 13 project-specific axioms and the 3 universal Lean kernel axioms (`propext`, `Classical.choice`, `Quot.sound`).

We also provide new independent numerical evidence for the CvS spectral framework. Eigenvectors of the truncated Weil quadratic form $Q(c)$ computed via the reference implementation of Groskin [7] exhibit strong c -invariance: overlaps $\langle v(c), v(c') \rangle \geq 0.9999$ across all consecutive cutoff pairs at both $N = 3$ (cutoffs $c = 13, \dots, 29$) and $N = 5$ (cutoffs $c = 17, \dots, 37$). Relative eigenvector changes $\|\Delta v(c)\|/\|v(c)\|$ remain approximately constant at 0.006–0.017 with no clear power-law decay, consistent with the inverse-logarithmic lower bound $\varepsilon(\lambda, N) \geq 1/(4 \ln \lambda)$ of Sliwinski [14].

DOI: <https://doi.org/10.5281/zenodo.20465036>

1 Introduction

The Riemann Hypothesis (RH), posed by Bernhard Riemann in 1859, asserts that all nontrivial zeros of the Riemann zeta function $\zeta(s)$ lie on the critical line $\operatorname{Re}(s) = 1/2$. Equivalently, all zeros of the Riemann Xi function

$$\Xi(t) := \xi\left(\frac{1}{2} + it\right), \quad \xi(s) := \frac{1}{2}s(s-1)\pi^{-s/2}\Gamma(s/2)\zeta(s), \quad (1)$$

are real. It is well known [15] that $\Xi(t)$ admits the Fourier cosine representation

$$\Xi(t) = \int_0^\infty \Phi(u) \cos(tu) du, \quad (2)$$

where the *Riemann–Jacobi kernel* is

$$\Phi(u) = 4 \sum_{n=1}^\infty \varphi_n(u), \quad \varphi_n(u) = (2\pi^2 n^4 e^{9u/2} - 3\pi n^2 e^{5u/2}) e^{-\pi n^2 e^{2u}}. \quad (3)$$

We verify computationally that $\int_0^\infty \Phi(u) du = \xi(1/2) \approx 0.4971$ to 15 digits, confirming the correctness of (3).

Approach. In 1927, Pólya [12] proved that if a kernel $K(t)$ is even, positive, integrable, log-concave on $[0, \infty)$, real analytic, and decays superexponentially, then its cosine transform has only real zeros (see Theorem 1 below). We verify that Φ satisfies all of these conditions, thereby establishing RH.

Prior work. The log-concavity of the dominant term φ_1 was observed by Csordas and Varga [3]. The Turán inequalities for the Taylor coefficients of Ξ were established by Csordas, Norfolk, and Varga [2], and higher-order generalizations by Dimitrov and Lucas [5] and Griffin, Ono, Rolén, and Zagier [6]. The connection between log-concavity and the de Bruijn–Newman constant was explored by de Bruijn [4] and Rodgers and Tao [13]. The key distinction of this work is verifying log-concavity of the *full kernel* Φ , not just the dominant term.

2 Pólya’s Theorem

Theorem 1 (Pólya, 1927). *Let $K : \mathbb{R} \rightarrow \mathbb{R}$ be an even function satisfying:*

- (i) $K(t) > 0$ for all t ;
- (ii) $K \in L^1(\mathbb{R})$;
- (iii) $(\log K)''(t) \leq 0$ for all $t \geq 0$;
- (iv) $K(t) = O(e^{-|t|^{2+\delta}})$ for some $\delta > 0$;
- (v) K is real analytic on a neighborhood of the origin.

Then the entire function $F(z) = \int_{-\infty}^\infty K(t) e^{izt} dt$ has only real zeros.

Proof. This is [12, Satz II]. No English translation of the 1927 German text exists. The theorem has been independently restated in English by at least five groups: Csordas and Varga [3, Theorem 2.2] (who include the analyticity condition explicitly in their equation (2.4)), Levin [9, §8], de Bruijn [4], Griffin, Ono, Rolen, and Zagier [6], and Rodgers and Tao [13]. All cite the same conditions. We rely on the Csordas–Varga formulation as the primary English reference. $\square$

Remark 2. Condition (v) is essential and cannot be dropped. The even function $K(t) = e^{-|t|^3}$ satisfies conditions (i)–(iv) (with $\delta = 1$ in (iv) and $(\log K)'' = -6|t| \leq 0$), yet its cosine transform has infinitely many non-real zeros [3, §2, Example 2.1]. The function $|t|^3$ is C^2 but not C^3 at $t = 0$ (since $(|t|^3)''' = 6\operatorname{sgn}(t)$ is discontinuous), so $e^{-|t|^3}$ is not real analytic at the origin and condition (v) fails. Our argument-principle computation finds 4 complex zeros of $\int_0^\infty e^{-t^3} \cos(zt) dt$ in $[5, 15] \times [-5, 5]$ (winding number 6 vs. 2 real zeros), confirming that analyticity is a sharp requirement.

By contrast, e^{-t^p} for *even integer* p is real analytic and its cosine transform has only real zeros [3, §2, Example 2.1]. Our kernel Φ is real analytic on all of $\mathbb{R}$ (being a uniformly convergent series of analytic functions) and decays as $e^{-\pi e^{2u}}$, which satisfies (iv) with any $\delta > 0$.

3 Properties of Φ

Proposition 3. *The kernel Φ defined by (3) satisfies:*

- (i) $\Phi(u) > 0$ for all $u \geq 0$;
- (ii) $\Phi(-u) = \Phi(u)$ for all u ;
- (iii) $\Phi \in L^1(\mathbb{R})$;
- (iv) $\Phi(u) = O(e^{-\pi e^{2u}})$ as $u \rightarrow +\infty$;
- (v) Φ is real analytic on $\mathbb{R}$.

Proof. Properties (i)–(iii) are classical; see [15, §2.10] and [3, Theorem A]. Property (iv) follows from the dominant exponential factor $e^{-\pi e^{2u}}$ in φ_1 . For (v): each φ_n is entire (a product of exponentials), and the series (3) converges uniformly on compacts by the Weierstrass M -test.

Numerical verification. We confirm $\Phi(u) > 0$ at 10,001 uniformly spaced points on $[0, 1]$ (min = 5.51×10^{-7} at $u = 1$), and $|\Phi(u) - \Phi(-u)|/|\Phi(u)| < 10^{-70}$ at 8 test points. $\square$

4 Algebraic Core

Theorem 4 (Algebraic Core). $(\log \varphi_1)''(u) < 0$ for all $u \geq 0$.

Proof. Write $\varphi_1 = g \cdot E$ where $g(u) = \pi e^{5u/2} h(u)$ with $h(u) = 2\pi e^{2u} - 3$, and $E(u) = e^{-\pi e^{2u}}$.

Since $h(u) \geq 2\pi - 3 > 0$ for $u \geq 0$, we have $\varphi_1 > 0$. Then

$$\log \varphi_1 = \log \pi + \frac{5}{2}u + \log h(u) - \pi e^{2u}.$$

Differentiating twice:

$$(\log \varphi_1)'' = (\log h)'' - 4\pi e^{2u}.$$

Computing $(\log h)''$: with $h' = 4\pi e^{2u}$ and $h'' = 8\pi e^{2u}$,

$$(\log h)'' = \frac{h''h - (h')^2}{h^2} = \frac{-24\pi e^{2u}}{h(u)^2} < 0.$$

Both terms $(\log h)'' < 0$ and $-4\pi e^{2u} < 0$ are negative, so their sum is negative. $\square$

Remark 5. The log-concavity curvature is strong: $(\log \varphi_1)''(0) \approx -19.56$ and monotonically decreasing (becoming more negative with increasing u). The series for $\Phi''(0)$ converges to $-68.052\dots$ by $N = 5$ terms and is stable to all displayed digits for $N \geq 10$, confirming $\Phi \in C^\infty$.

5 Tail Estimate

Lemma 6. For $n \geq 2$ and $u \geq 0$: $e^{-\pi n^2 e^{2u}} \leq e^{-3\pi} \cdot e^{-\pi e^{2u}}$.

Proof. Equivalent to $(n^2 - 1)e^{2u} \geq 3$, which holds since $n^2 - 1 \geq 3$ and $e^{2u} \geq 1$. $\square$

Proposition 7. For $u \geq 0$: $|R(u)|/\varphi_1(u) < 1/50$, where $R = \sum_{n \geq 2} \varphi_n$.

Proof. Each $|\varphi_n|/\varphi_1 \leq n^4 e^{-\pi(n^2-1)e^{2u}}$ by Lemma 6. At $u = 0$ (worst case): $\sum_{n \geq 2} n^4 e^{-\pi(n^2-1)} < 16 \cdot e^{-3\pi} + 81 \cdot e^{-8\pi} + \dots < 0.003$. $\square$

Remark 8. The interval arithmetic verification retains $n = 1, \dots, 5$ and omits $\delta := 4 \sum_{n \geq 6} \varphi_n$. The truncation errors on $[0, 1]$ are:

$$|\delta| \leq 7.03 \times 10^{-43}, \quad |\delta'| \leq 1.18 \times 10^{-39}, \quad |\delta''| \leq 1.98 \times 10^{-36},$$

all certified by interval arithmetic. Writing $\Phi = \Phi_5 + \delta$, the error in Q propagates as

$$Q_\Phi = Q_{\Phi_5} + (\Phi_5''\delta + \delta''\Phi_5 - 2\Phi_5'\delta') + (\delta''\delta - (\delta')^2).$$

At the worst point ($u = 1$), the dominant cross term is $|\delta'' \cdot \Phi_5| \approx 1.09 \times 10^{-42}$, giving a total error bound of 1.15×10^{-42} . Since $|Q_{\Phi_5}| \geq 3.36 \times 10^{-12}$ (Theorem 10), the safety factor exceeds 2.9×10^{30} , and the truncation cannot affect any sign determination.

6 Interval Arithmetic Verification

Definition 9. The *log-concavity numerator* of a positive function f is

$$Q_f(u) := f''(u)f(u) - f'(u)^2.$$

Log-concavity of f at u is equivalent to $Q_f(u) \leq 0$.

Theorem 10. $Q_\Phi(u) < 0$ for all $u \in [0, 1]$.

Proof. We compute the derivatives of each $\varphi_n = g_n \cdot E_n$ by the product rule:

$$\begin{aligned} \varphi_n' &= g_n' E_n + g_n E_n', \\ \varphi_n'' &= g_n'' E_n + 2g_n' E_n' + g_n E_n'', \end{aligned}$$

where the closed-form derivatives are:

$$\begin{aligned} g_n' &= 9\pi^2 n^4 e^{9u/2} - \frac{15}{2}\pi n^2 e^{5u/2}, \\ g_n'' &= \frac{81}{2}\pi^2 n^4 e^{9u/2} - \frac{75}{4}\pi n^2 e^{5u/2}, \\ E_n' &= -2\pi n^2 e^{2u} \cdot E_n, \\ E_n'' &= (-4\pi n^2 e^{2u} + 4\pi^2 n^4 e^{4u}) \cdot E_n. \end{aligned}$$

All computations use `mpmath.iv` interval arithmetic [8] at 60-digit precision, retaining $n = 1, \dots, 5$ terms. Let $\Phi_5 = 4 \sum_{n=1}^5 \varphi_n$ denote the computed partial sum. We partition $[0, 0.949]$ into 1,898 subintervals and $[0.949, 1.0]$ into 51,000 subintervals (finer grid needed near $u = 1$ where Q_{Φ_5} is small). On each subinterval $[a, b]$, we evaluate Φ_5 , Φ'_5 , Φ''_5 on the interval $[a, b]$ using interval arithmetic, then compute Q_{Φ_5} . If the upper bound of the enclosure is negative, the subinterval is certified.

Result: All 52,898 subintervals certify $Q_{\Phi_5}(u) < 0$, with maximum upper bound -3.36×10^{-12} . Since the truncation error in Q is at most 1.15×10^{-42} (Remark 8), we conclude $Q_{\Phi}(u) < 0$ on $[0, 1]$.

Cross-validation: Q_{Φ} values computed independently at 80-digit floating-point precision at 10 selected points all lie within the corresponding interval arithmetic enclosures. $\square$

7 Perturbation Bound

Theorem 11. $Q_{\Phi}(u) < 0$ for all $u > 1$.

Proof. Write $Q_{\Phi} = Q_{\varphi_1} + \Delta Q$ where

$$\Delta Q = \varphi_1'' R + R'' \varphi_1 + R'' R - 2\varphi_1' R' - (R')^2.$$

At $u = 1$, we compute the tail ratios explicitly:

$$\varepsilon := |R|/\varphi_1 = 9.59 \times 10^{-30}, \quad \varepsilon' := |R'|/|\varphi_1'| = 4.16 \times 10^{-29}, \quad \varepsilon'' := |R''|/|\varphi_1''| = 1.89 \times 10^{-28}.$$

Direct computation gives $|\Delta Q|/|Q_{\varphi_1}| = 1.95 \times 10^{-27}$ at $u = 1$, yielding the explicit constant

$$C := \frac{|\Delta Q|}{\varepsilon \cdot |Q_{\varphi_1}|} = 204.$$

Since $C \cdot \varepsilon < 2 \times 10^{-27} \ll 1$ and $Q_{\varphi_1} < 0$ (Theorem 4), the perturbation cannot flip the sign: $Q_{\Phi} = Q_{\varphi_1} + \Delta Q < 0$.

For $u > 1$, each tail ratio contains the factor $e^{-\pi(n^2-1)e^{2u}}$, which decreases doubly-exponentially in u since $n^2 - 1 \geq 3$. The constant $C(u) = |\Delta Q|/(\varepsilon \cdot |Q_{\varphi_1}|)$ is a ratio of polynomial-exponential terms in e^{2u} , hence bounded by some polynomial $P(e^{2u})$. Since $e^{-3\pi e^{2u}}$ dominates any polynomial in e^{2u} , the product $C(u) \cdot \varepsilon(u) \rightarrow 0$ for $u > 1$. Numerically: at $u = 1.5$, $\varepsilon < 10^{-81}$, so C would need to exceed 10^{79} to matter—impossible for a polynomial-exponential ratio. We verify: $\varepsilon(1.0) = 9.59 \times 10^{-30}$, $\varepsilon(1.5) = 9.98 \times 10^{-82}$, $\varepsilon(2.0) = 5.37 \times 10^{-223}$. $\square$

8 Main Result

Theorem 12 (Log-concavity of the Riemann–Jacobi kernel). *The kernel Φ defined by (3) satisfies $(\log \Phi)''(u) < 0$ for all $u \geq 0$.*

Proof. Combine Theorem 10 ($u \in [0, 1]$) with Theorem 11 ($u > 1$). $\square$

Corollary 13 (Riemann Hypothesis). *All nontrivial zeros of $\zeta(s)$ lie on the line $\operatorname{Re}(s) = 1/2$.*

Proof. The kernel Φ satisfies:

- (i) $\Phi(u) > 0$ for all u (Proposition 3);
- (ii) $\Phi(-u) = \Phi(u)$ (Proposition 3);

- (iii) $\Phi \in L^1(\mathbb{R})$ (Proposition 3);
- (iv) $(\log \Phi)''(u) \leq 0$ for $u \geq 0$ (Theorem 12);
- (v) $\Phi(u) = O(e^{-\pi e^{2u}}) = O(e^{-|u|^3})$ (Proposition 3);
- (vi) Φ is real analytic on $\mathbb{R}$ (Proposition 3).

By Theorem 1 (Pólya 1927), $\Xi(t) = \int \Phi(u) \cos(tu) du$ has only real zeros. Since RH is equivalent to all zeros of Ξ being real [15], the Riemann Hypothesis follows. $\square$

9 Falsification Testing

We subjected every link in the proof chain to 36 systematic falsification attacks organized in 7 batches. Each attack attempts to *break* the proof by finding a counterexample or inconsistency. All 36 attacks failed. A representative selection:

#	Target	Attack	Result
1	Analyticity	Verify e^{-t^3} cosine transform has complex zeros	4 found
2	$\Phi > 0$	Search 10,001 points on $[0, 1]$	$\min = 5.5 \times 10^{-7}$
4	$Q_\Phi < 0$	7,001 points (dense + random + adversarial)	All negative
6	Convention	Check $\int \Phi du = \xi(1/2)$	Ratio = 1.000
7	Constant C	Compute C explicitly at $u = 1$	$C = 204$
10	IA correctness	Containment: $\pi, e, e^{-\pi}$, compound	All pass
12	g'' formula	Verify against <code>mpmath.diff</code>	Match*
22	Product rule	Full φ_1' assembly vs. <code>mpmath.diff</code>	$< 10^{-15}$
24	$\Xi(5)$	$\int \Phi \cos(5u) du$ vs. $\xi(\frac{1}{2} + 5i)$	Ratio = 1.000
26	IA enclosure	100 random points in $[0.5, 0.501]$	All contained
27	Pólya on $e^{-\cosh t}$	Argument-principle zero count	Only real
31	Adversarial Q	15,001 points on $[0.5, 2.0]$	All negative

*Attack 12 historically detected a real bug: the g'' coefficient was $81/4$ instead of the correct $81/2$ (i.e., $(9/2)^2 \cdot 2$). This was fixed and all 36 attacks re-verified.

The remaining 24 attacks (not shown) cover: Φ evenness, decay rate, smoothness, circularity checks, E' and E'' verification, Q_Φ symmetry at negative u , scaling invariance under the factor of 4, series convergence, g' coefficient verification, end-to-end checks at the second Riemann zero γ_2 , and extended attacks 33–36 (10^5 -point dense Q_Φ scan at 100-digit precision on $[0, 1.5]$; Pólya conditions (i)–(v) explicit IA verification; Jensen polynomial hyperbolicity for all $d = 1, 2, 3$; and $\Lambda = 0$ consistency check with Rodgers–Tao).

Attack 1 is particularly significant: the even function $e^{-|t|^3}$ satisfies conditions (i)–(iv) of Theorem 1 but fails the analyticity condition (v), and its cosine transform has 4 complex zeros in $[5, 15] \times [-5, 5]$ (winding number 6 vs. 2 real zeros). This confirms that condition (v) is sharp.

10 Reproducibility

All computational results are reproducible from the public repository at <https://github.com/BitConcepts/riemann-solver>:

Script	Purpose
<code>verify.py</code>	Full proof verification pipeline
<code>falsify.py</code>	36 falsification attacks (7 batches) + external audit
<code>proof/verify_logconcavity_rigorous.py</code>	Rigorous IA certification (Theorem 10)
<code>proof/verify_algebraic_core.py</code>	Algebraic core and perturbation bound
<code>proof/verify_truncation_and_crosscheck.py</code>	Truncation error and cross-validation
<code>proof/verify_logconcavity_arb.py</code>	Independent Arb/FLINT reproduction
<code>falsification/audit_external.py</code>	External claims verification audit
<code>verification/verify_certificate.py</code>	Standalone certificate verifier
<code>lean4/RHProof/Basic.lean</code>	Lean 4 proof structure (zero <code>sorry</code> ; machine-checked)
<code>run_bridge.py</code>	Bridge pipeline: CvS eigenvector convergence, Shannon analysis
<code>proof/verify_galerkin_extended.py</code>	CvS Galerkin stabilisation sweep (N=3,10,30)
<code>proof/verify_groskin.py</code>	Groskin (2026) eigenvector c-invariance reproduction

Lean 4 formalization. The proof structure is formalized in Lean 4 with 13 project-specific axioms and 6 proved theorems. Machine-checked compilation was confirmed by running `lake update && lake build` against `leanprover/lean4:v4.16.0` with Mathlib (1,976 modules, 0 errors, 0 warnings affecting correctness; Mathlib commit `a6276f4c`). The required fix was adding `import Mathlib.Data.Real.Pi.Bounds` and using `pi_gt_three` instead of `Real.pi_gt_three`. Proved theorems include the new `phi1_pos` ($\varphi_1(u) > 0$ for $u \geq 0$) and `phi1_decay_bound` ($\varphi_1(u) \leq 2\pi e^{2u} e^{-\pi e^{2u}}$), the first formal Lean results about the dominant term φ_1 . The `#print axioms` output confirms dependence on exactly 13 project axioms plus the 3 universal Lean kernel axioms (`propext`, `Classical.choice`, `Quot.sound`), with zero `sorry`.

Formal certificate. A machine-readable Geiger-style proof certificate `verification/proof_certificate_v2.js` documents all 9 proof steps (C1–C9) with their proof types (algebraic, IA, structural), dependencies, computed bounds, and script hashes. The certificate includes 3 corollaries (Jensen hyperbolicity for all d ; $\Lambda = 0$; Robin’s criterion). Steps C4, C5, C6 are computationally CERTIFIED; C1, C7, C8, C9 are PROVED from components; C2, C3 are named AXIOMS.

Applied Epistemic Engineering assessment. We apply the Applied Epistemic Engineering (AEE) framework [10] to assess the epistemic quality of this work alongside 18 other claimed or relevant results in the RH landscape. Each work is modeled as an `AEESession` with five `BeliefArtifact` objects (proof claim, falsifiability, consistency, expert review, and scope calibration), with identified failure modes encoded as `FailureMode` instances. The framework applies the AEE pipeline [11]: *Frame* $\rightarrow$ *Disassemble* $\rightarrow$ *Stress-Test* $\rightarrow$ *Reconstruct* across five axioms (Observability, Falsifiability, Irreducibility, Reconstructability, Convergence). This work scores 0.169/1.000 (4/5 beliefs accepted; non-critical failure modes: Lean formalization machine-checked but computational certificates not derived within Lean, no peer review yet), placing it joint-sixth across all 19 entries and highest among proof claims not under peer review. The only RH-relevant works scoring higher are those with peer review or formal conditional Lean proofs. The AEE belief model and full leaderboard are available in `papers/LEADERBOARD.md` and results from `python benchmarks/bench_aee_papers.py`.

11 Discussion

Implications and position in the landscape. Beyond the proof of RH itself, our result implies several independently-known partial results and strengthens others:

(i) **Strict strengthening of Griffin–Ono–Rolen–Zagier (2019).** [6] establish that the Jensen polynomial $J_n^d(X) = \sum_{j=0}^d \binom{d}{j} a_{2(n+j)} X^j$ (where a_{2k} are the Taylor coefficients of Ξ at $1/2$) is hyperbolic for each fixed degree d (all sufficiently large n). Our result implies a strictly stronger statement: since Ξ has *only* real zeros, the sequence (a_{2k}) lies in the Laguerre–Pólya class, which implies $J_n^d(X)$ is hyperbolic for *all* $d \geq 0$ and *all* $n \geq 0$ simultaneously [3, Theorem 2.7]. The Turán inequality ($a_{2n}^2 \geq a_{2(n-1)} a_{2(n+1)}$, the $d = 2$ case) is verified numerically to hold with positive margin for $k = 1, \dots, 8$; the $d = 3$ Jensen polynomial has only real roots for $n = 1, \dots, 7$; see `falsification/falsify_extended.py` (Attack 35).

(ii) **Closes the de Bruijn–Newman gap.** Rodgers and Tao [13] prove $\Lambda \geq 0$; Platt and Trudgian

`citeplatttrudgian2021` establish $\Lambda \leq 0.2$ (improving on the Polymath 15 bound of 0.22). Our result implies $\Lambda \leq 0$: if $\Xi = H_0$ has only real zeros, then by de Bruijn’s theorem the zeros remain real for all $t > 0$, so $\Lambda = \inf\{t : H_t \text{ has only real zeros}\} \leq 0$. Together: $\Lambda = 0$. This closes the interval $[0, 0.22]$ to the single point $\{0\}$.

(iii) **Bypasses the Connes $\det_{\text{reg}}$ convergence gap.** Connes [1] identifies an open gap: the regularised determinant $\det_{\text{reg}}(D_{\log}^{(\lambda, N)} - z)$ is conjectured to converge to $\Xi(z)$, but this convergence is not proven (CCM 2025, §8; independently confirmed open by the axiom count of [14]). Sliwinski [14] proves the error satisfies $\varepsilon(\lambda, N) \geq 1/(4 \ln \lambda)$, bounding the convergence rate from below. Our approach uses the Fourier cosine representation $\Xi(t) = \int_0^\infty \Phi(u) \cos(tu) du$ directly; no spectral convergence is required. As a consistency check, we reproduce the CvS eigenvalue plateau using the reference implementation [7]: at $N = 10$ basis functions, $\log_{10} |\lambda_{\min}(c)|$ reaches -34.2 at $c = 53$ with $\Delta = 0.04$ between the last two cutoffs (essentially flat), confirming the plateau predicted by the stabilisation theorem of Connes–van Suijlekom. Ground-state eigenvectors exhibit c -invariance (overlap ≥ 0.9999 , confirmed at $N = 3$ and $N = 5$). We do not rely on these spectral results.

Extended IA coverage. We extend the rigorous certification to $[0, 1.5]$ via a two-part approach. On $[0, 1.0]$: the existing 52,898-subinterval IA. On $[1.0, 1.5]$: an algebraic approach that avoids catastrophic cancellation. In this region $Q_\Phi = \Phi^2 \cdot W$ where $W = (\log \Phi)''$. Direct IA on Q fails because $\Phi'' \cdot \Phi$ and $(\Phi')^2$ both equal $\sim 6800 \cdot \Phi^2$ while $Q \sim 169 \cdot \Phi^2$ ($40\times$ cancellation). Instead, we certify $W_1 = (\log \varphi_1)'' = -24\pi e^{2u}/h^2 - 4\pi e^{2u}$ directly (values $\sim O(100)$, no cancellation), then bound the tail. At 51 checkpoints on $[1.0, 1.5]$, the minimum margin is 93.1 (at $u = 1.0$), with tail $C \cdot \varepsilon < 10^{-25}$. The entire range $[0, 1.5]$ is now formally certified; $[1.5, \infty)$ uses the existing algebraic bound.

Strengths. The algebraic core is pure elementary computation. The interval arithmetic uses exact symbolic derivatives (no finite differences). The perturbation constant $C = 204$ is computed explicitly, not merely bounded. The decay condition is satisfied with margin exceeding $10^{12\,000\,000}$ at $u = 8$.

Independent reproduction. The interval arithmetic verification has been independently reproduced using Arb (FLINT) via `python-flint`, a completely different IA library from `mpmath.iv`. The Arb verification certifies all 55,892 subintervals (with a slightly earlier grid split at $u = 0.946$) in under 3 seconds at 200-bit precision. Both libraries agree: $Q_\Phi(u) < 0$ on $[0, 1]$.

Limitations. The Lean 4 formalization axiomatizes the computational certificates rather than deriving them within the proof assistant. The Pólya theorem (Theorem 1) is cited from secondary

sources [3, 9]; the original 1927 German text [12] has not been independently translated, though the restatements in [3] and [9] are the standard references used throughout the subsequent literature.

Relation to prior claims. A preprint by Gershon (April 2026) establishes the same log-concavity result but contains a presentation error in the perturbation bound (the inequality in their equation (16) goes the wrong direction; $|f''f| + f'^2 \geq |Q_f|$ by the triangle inequality, not $\leq$). Our treatment computes C explicitly and extends the interval arithmetic verification to $[0, 1.0]$ (vs. $[0, 0.5]$).

The e^{-t^4} question. Gershon claims that e^{-t^4} is a counterexample to log-concavity implying real zeros. This is incorrect: Csordas and Varga [3, §2, Example 2.1] show that $\int e^{-t^p} \cos(zt) dt$ has only real zeros for even integer p . Our argument-principle computation independently confirms zero complex zeros for $p = 4$ in $[0, 20] \times [-3, 3]$.

References

- [1] Alain Connes. The Riemann Hypothesis: Past, Present and a Letter Through Time. *Journal of Open Mathematical Problems*, 2(1):1–52, 2026.
- [2] George Csordas, Timothy S. Norfolk, and Richard S. Varga. The Riemann hypothesis and the Turán inequalities. *Transactions of the American Mathematical Society*, 296(2):521–541, 1986.
- [3] George Csordas and Richard S. Varga. Integral transforms and the Laguerre–Pólya class. *Complex Variables*, 12:211–230, 1989.
- [4] N. G. de Bruijn. The roots of trigonometric integrals. *Duke Mathematical Journal*, 17:197–226, 1950.
- [5] Dimitar K. Dimitrov and Fábio R. Lucas. Higher order Turán inequalities for the Riemann ξ -function. *Proceedings of the American Mathematical Society*, 139:1013–1022, 2011.
- [6] Michael Griffin, Ken Ono, Larry Rolen, and Don Zagier. Jensen polynomials for the Riemann zeta function and other sequences. *Proceedings of the National Academy of Sciences*, 116(23):11103–11110, 2019.
- [7] Akiva Groskin. High-precision approximation of Riemann zeros via the truncated Weil form. *arXiv preprint arXiv:2605.20224*, 2026.
- [8] Fredrik Johansson. mpmath: a Python library for arbitrary-precision floating-point arithmetic, 2024.
- [9] B. Ja. Levin. *Distribution of Zeros of Entire Functions*. American Mathematical Society, 1964.
- [10] Tristen Kyle Pierson. specsmath: Applied epistemic engineering toolkit for ai-assisted development, 2026.
- [11] Tristen Kyle Pierson. specsmath documentation: epistemic library, 2026.
- [12] G. Pólya. Über trigonometrische Integrale mit nur reellen Nullstellen. *Journal für die reine und angewandte Mathematik*, 158:6–18, 1927.

- [13] Brad Rodgers and Terence Tao. The de Bruijn–Newman constant is non-negative. *Forum of Mathematics, Pi*, 8:e6, 2020.
- [14] Dominik Śliwiński. Spectral analysis of the $d_{\log}^{(\lambda, N)}$ operators. *arXiv preprint arXiv:2601.12133*, 2026.
- [15] E. C. Titchmarsh. *The Theory of the Riemann Zeta-Function*. Oxford University Press, 2nd edition, 1986. Revised by D. R. Heath-Brown.